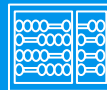




recod.ai
reasoning for complex data



Ensemble Learning

Machine Learning

Prof. Sandra Avila

Institute of Computing (IC/Unicamp)


MC886/MO444, November 22, 2022



**Guess how many
jelly beans are in
the jar**


WINNER RECEIVES A **FREE** MEAL
VOUCHER!

PUT NAME, EMAIL ADDRESS AND GUESS ON PAPER



Francis Galton
(1822-1909)

Animal's weight?



~800 people
542 kg

543 kg



Wisdom of the Crowd



Ensemble Learning

- Multiple learning algorithms **to obtain better predictive performance** than could be obtained from any learning algorithms individually.


Voting Classifiers

~80%

Logistic
Regression



Neural
Network




Random
Forest



Other ...




Diverse
Predictors



Voting Classifiers

Hard/Soft voting classifier



Voting Classifiers

- Voting classifier **often achieves a higher accuracy than the best classifier** in the ensemble.
- Even if each classifier is a **weak learner**, the ensemble can still be a **strong learner**, provided there are a sufficient number of weak learners and they are sufficiently diverse.

Voting Classifiers

- Ensemble methods work best when the predictors are as **independent** from one another as possible.
- One way to get diverse classifiers is to train them using **very different algorithms**: this increases the chance that they will make very different types of errors, improving the ensemble's accuracy.

Voting Classifiers

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
log_clf = LogisticRegression()
rnd_clf = RandomForestClassifier()
svm_clf = SVC()
voting_clf = VotingClassifier(
    estimators=[('lr', log_clf), ('rf', rnd_clf), ('svc', svm_clf)],
    voting='hard'
)
voting_clf.fit(X_train, y_train)
```

```
LogisticRegression 0.864
RandomForestClassifier 0.896
SVC 0.888
VotingClassifier 0.904
```

Today's Agenda

— — —


- Ensemble Methods
 - Bagging (and Pasting)
 - Boosting
 - Stacking

Bagging & Pasting

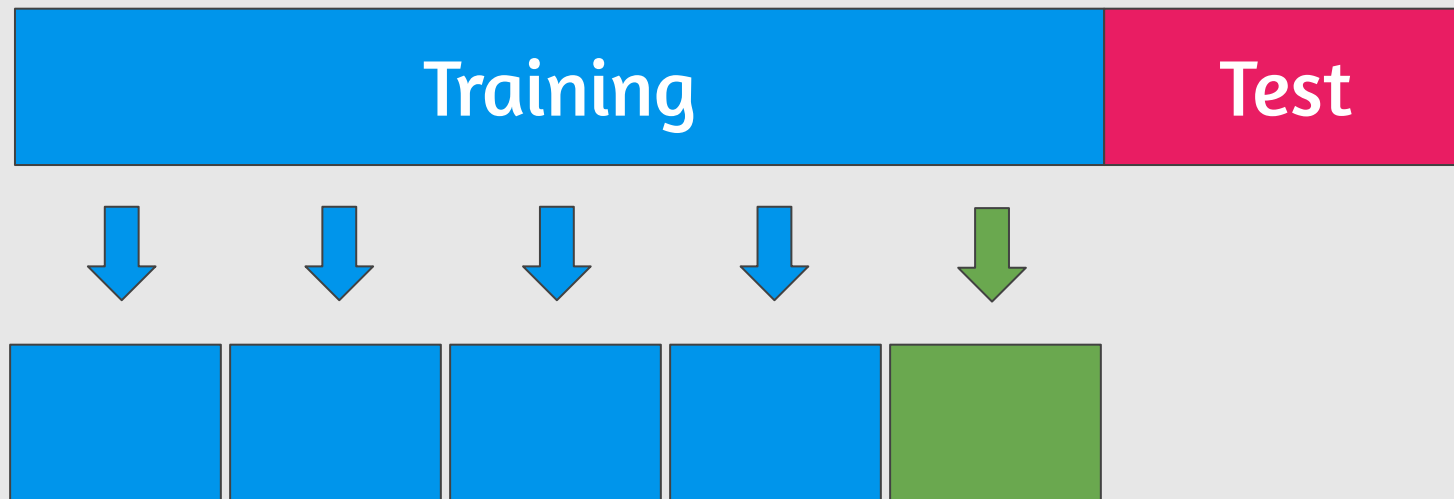
Bagging and Pasting

- Use **the same training algorithm** for every predictor, but to train them on **different random subsets** of the training set.
- **Bagging** (short for Bootstrap Aggregating): sampling is performed **with** replacement.
- **Pasting**: sampling is performed **without** replacement.

Bagging and Pasting



Bagging vs. Cross Validation



Training

Test

Cross
Validation

Training

Test

Cross
Validation
(one model)

Training

Test

Random subset

Random subset

Random subset

Random subset






Random subset



Bagging
(many models)

Training

Test



Bagging

Bagging and Pasting

- Once all predictors are trained, the ensemble can make a prediction for a new instance by simply **aggregating the predictions of all predictors**.
- **Bagging and Pasting scale very well.**

Bagging and Pasting

```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
bag_clf = BaggingClassifier(
    DecisionTreeClassifier(), n_estimators=500,
    max_samples=100, bootstrap=True, n_jobs=-1
)
bag_clf.fit(X_train, y_train)
y_pred = bag_clf.predict(X_test)
```

Today's Agenda

— — —

- Ensemble Methods
 - Bagging (and Pasting)
 - **Boosting**
 - Stacking

Boosting


Boosting

- The general idea of most boosting methods is **to train predictors sequentially**, each trying to correct its predecessor.
- Most popular: AdaBoost and Gradient Boost.


AdaBoost [Freund and Schapire, 1997]

- One way for a new predictor to correct its predecessor is to pay a bit **more attention** to the training instances that **the predecessor underfitted**.
- This results in new predictors focusing more and more on **the hard cases**.


AdaBoost




AdaBoost




AdaBoost




AdaBoost




AdaBoost




AdaBoost




AdaBoost



AdaBoost




**Weak
Classifier 1**




**Weights
Increased**



**Weak
Classifier 2**



**Weak
classifier 3**



**Final classifier is
linear combination of
weak classifiers**




AdaBoost

1. Assign every observation, x_i , an initial weight value, $w_i = \frac{1}{n}$, where n is the total number of observations.
2. Train a "weak" model. (most often a decision tree)
3. For each observation:
 - 3.1. If predicted incorrectly, w_i is increased
 - 3.2. If predicted correctly, w_i is decreased
4. Train a new weak model where observations with greater weights are given more priority.
5. Repeat steps 3 and 4 until observations perfectly predicted or a preset number of trees are trained.

Chris Albon

Gradient Boosting [Breiman, 1997]

- Instead of tweaking the instance weights at every iteration like AdaBoost does, this method fit the new predictor to the **residual errors** made by the previous predictor.
- Instead of training on a newly sample distribution, the weak learner **trains on the remaining errors**.



```
from sklearn.tree import DecisionTreeRegressor

tree_reg1 = DecisionTreeRegressor(max_depth=2)
tree_reg1.fit(X, y)

y2 = y - tree_reg1.predict(X)
tree_reg2 = DecisionTreeRegressor(max_depth=2)
tree_reg2.fit(X, y2)

y3 = y2 - tree_reg2.predict(X)
tree_reg3 = DecisionTreeRegressor(max_depth=2)
tree_reg3.fit(X, y3)

y_pred = sum(tree.predict(X_new) for tree in (tree_reg1, tree_reg2, tree_reg3))
```

```
from sklearn.ensemble import GradientBoostingRegressor

gbrt = GradientBoostingRegressor(max_depth=2, n_estimators=3, learning_rate=1.0)
gbrt.fit(X, y)
```


Gradient Boosting [Breiman, 1997]

1. Fit a simple linear regressor or decision tree on data
[call **x** as input and **y** as output]
2. Calculate error residuals. Actual target value, minus predicted target value
[**e1 = y - y_predicted1**]
3. Fit a new model on error residuals as target variable with same input variables
[call it **e1_predicted**]
4. Add the predicted residuals to the previous predictions
[**y_predicted2 = y_predicted1 + e1_predicted**]
5. Fit another model on residuals that is still left, i.e. [**e2 = y - y_predicted2**] and repeat steps 2 to 5 until it starts overfitting or the sum of residuals become constant.

Gradient Boosting [Breiman, 1997]

- XGboost [Chen and Guestrin, 2016]:

Extreme Gradient Boosting

<https://github.com/tqchen/xgboost>

It aims at being extremely fast, scalable and portable.

Today's Agenda

— — —

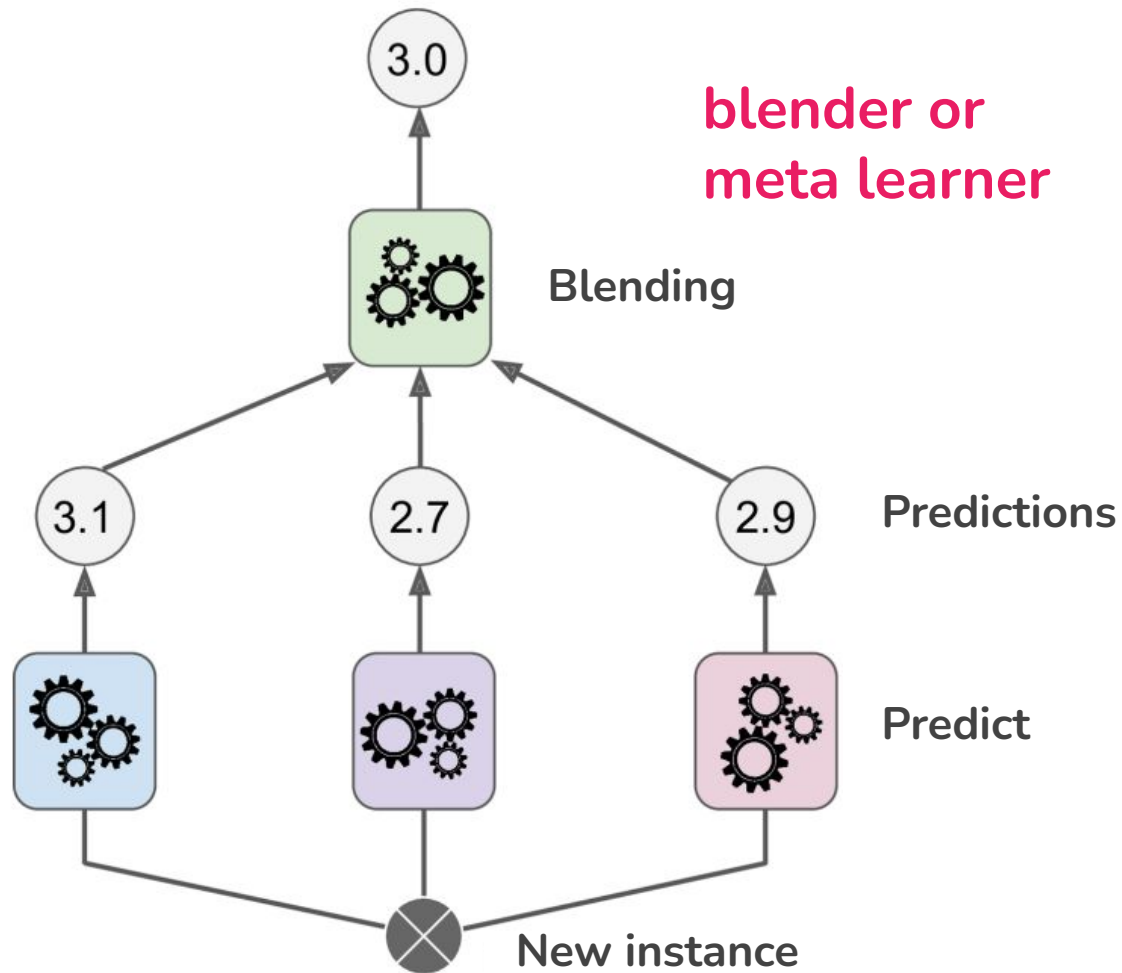
- Ensemble Methods
 - Bagging (and Pasting)
 - Boosting
 - **Stacking**

Stacking


Stacking [Wolpert, 1992]

- Stacking (short for Stacked Generalization)
- Instead of using trivial functions (such as hard voting) to aggregate the predictions of all predictors in an ensemble, we **train a model to perform this aggregation.**


Stacking



To train the
blender, a
common
approach is
to use a
hold-out set.



Multi-layer Stacking Ensemble



Stacking [Wolpert, 1992]

- ~~Scikit-Learn does not support stacking directly. =(~~
- <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.StackingClassifier.html>

References

— — —

Machine Learning Books

- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Chap. 6 & 7
- Pattern Recognition and Machine Learning, Chap. 14
- Pattern Classification, Chap 8 & 9 (Sec. 9.5)
- “Scikit Learn Ensemble Learning, Bootstrap Aggregating (Bagging) and Boosting” <https://youtu.be/X3Wbfb4M33w>